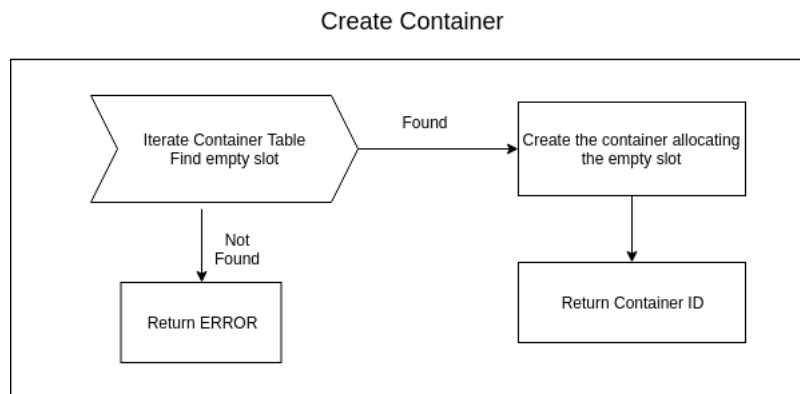


1 Container Manager

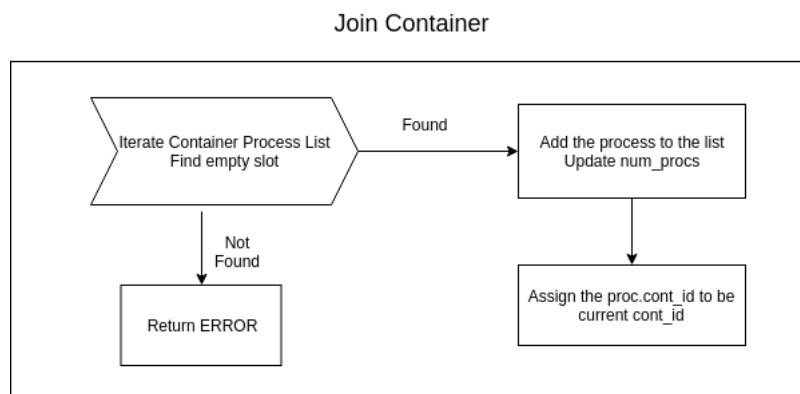
- This is implemented in the Kernel space.
- A container table stores all containers and relevant data structures for them. This table is analogous to the process table. The container data structure:

```
struct container_desc{  
    int allocated;           // whether the container is allocated  
    int procs[MAX_PROCS];    // List of processes in this container  
    int num_procs;           // Number of processes in this container  
    int last_run_id;         // PID of last scheduled process of the container  
};
```

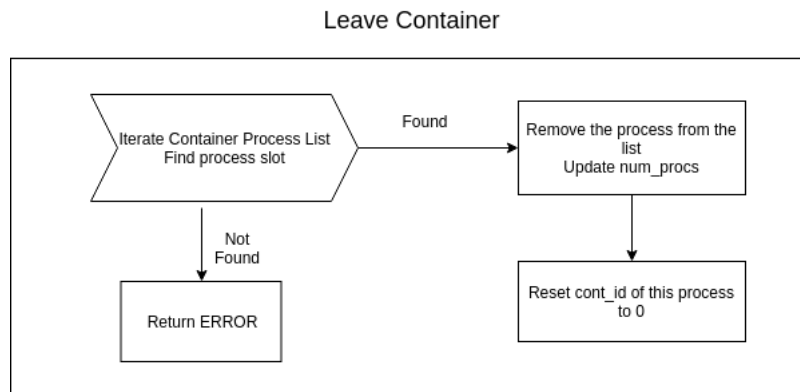
1.1 Container Creation



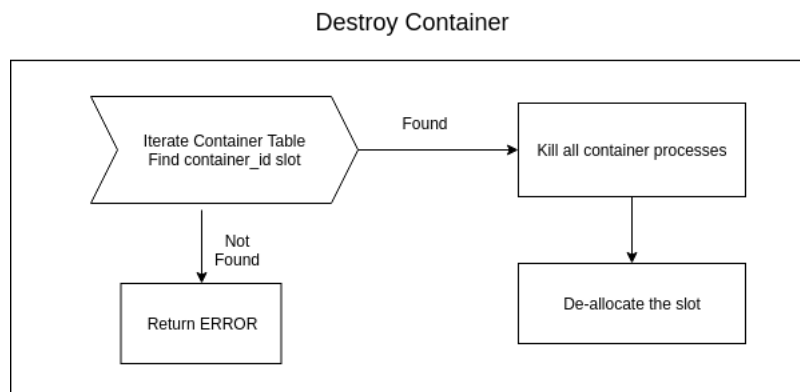
1.2 Joining a Container



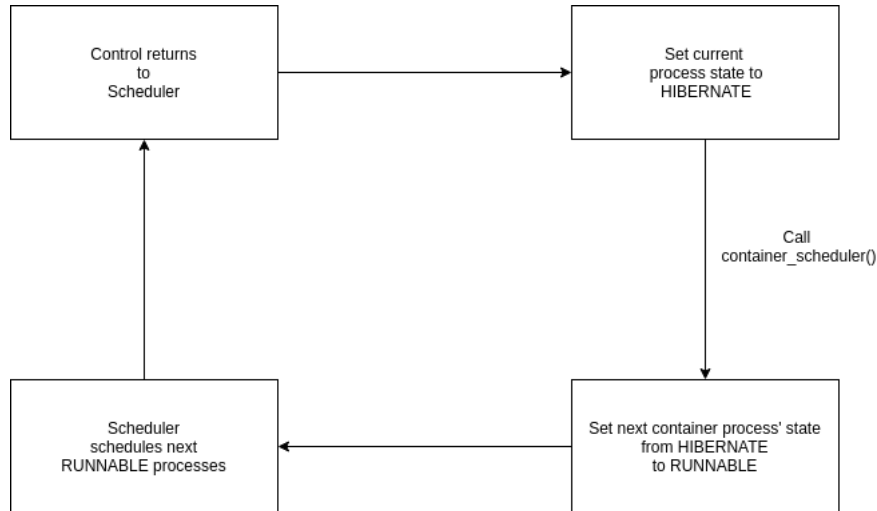
1.3 Leaving a Container



1.4 Destroying Container



2 Virtual Scheduler



We have implemented an intra-container round robin scheduling algorithm. Each container iterates over its contained processes in a round robin fashion and makes the next ready process available for scheduling.

1. The currently scheduled process returns the control back to scheduler.
2. If the returned process state is RUNNABLE, scheduler sets it to HIBERNATE, and *container_schedule()* is called with the current container id.
3. *container_schedule()* sets the state of next container process which was on HIBERNATE to RUNNABLE.
4. Control returns to scheduler which schedules the next RUNNABLE process.

2.1 ps syscall


To implement ps syscall respecting container behaviour, we modified the standard ps syscall implemented in assignment 1 to check for container id of all the active processes and then print only those processes whose container id was same as current process.

3 Virtual File system


We have an implicit table which produces a mapping from a given file name and the container ID of the process to give a unique **cfid**.

cfid	File Name	Container ID
2	my_file	1
5	my_file	2
8	file_3	1

3.1 ls syscall



3.2 open syscall



3.3 write syscall

